

CASCADE — Technical Design Document

Change Analysis for Software Cascade using Dependency Evidence

> *"Understand every consequence before every merge."*

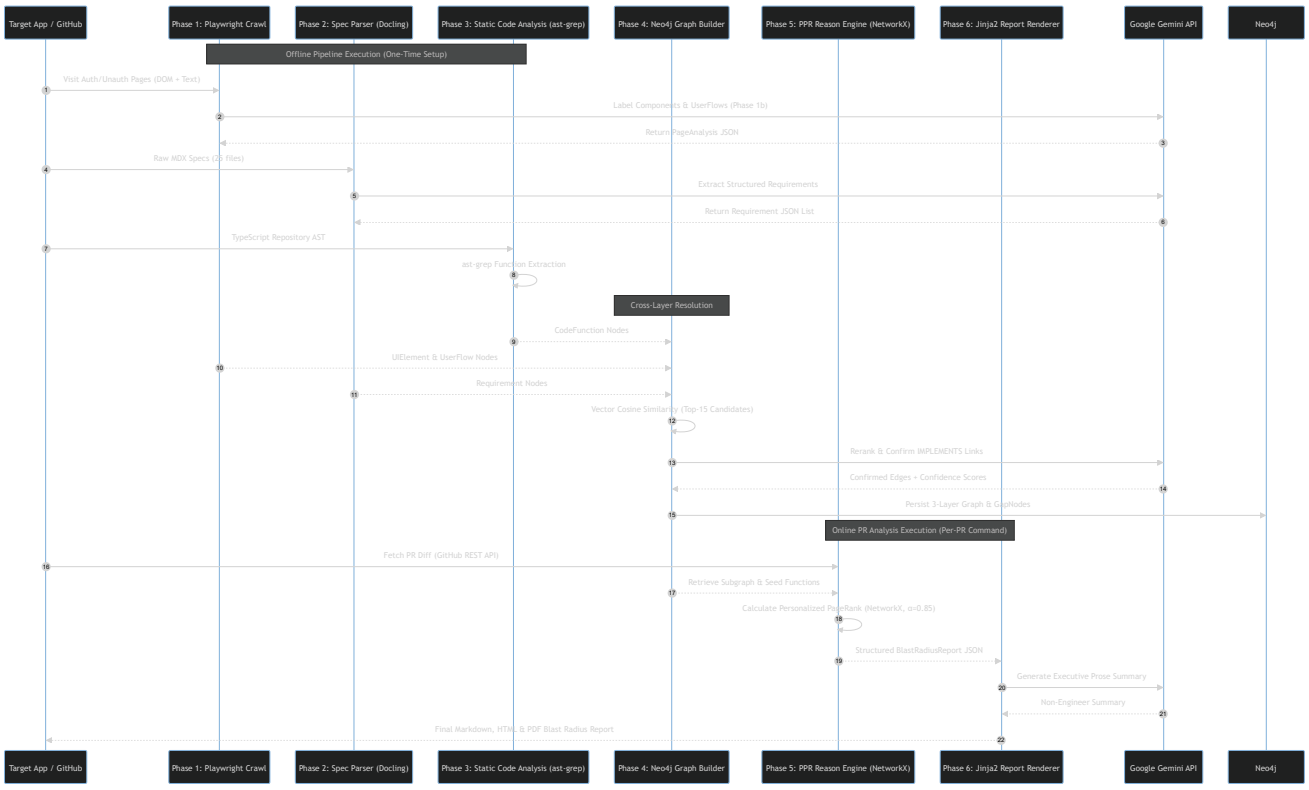
Candidate: Harsh Vardhan (*B.Tech CSE - Data Science & Analytics, IIIT Nagpur*)
Role: AI Engineer — Testsigma
Target Application: app.documenso.com / github.com/documenso/documenso
PR Analyzed: #3046 (*merge SHA* `8f683932`)
Pipeline Version: CASCADE v4 (6-Phase Hybrid Architecture)

---

1. Agent Decomposition

Stages of Reasoning, Boundaries & Determinism vs. LLM Trade-Offs

The Blast Radius Analyzer translates code changes into product-level impact. To avoid building a "chain of prompts dressed up as an agent," the pipeline explicitly separates **deterministic computation** (static AST parsing, graph linear algebra) from **LLM semantic reasoning** (document parsing, cross-domain reranking, prose synthesis).



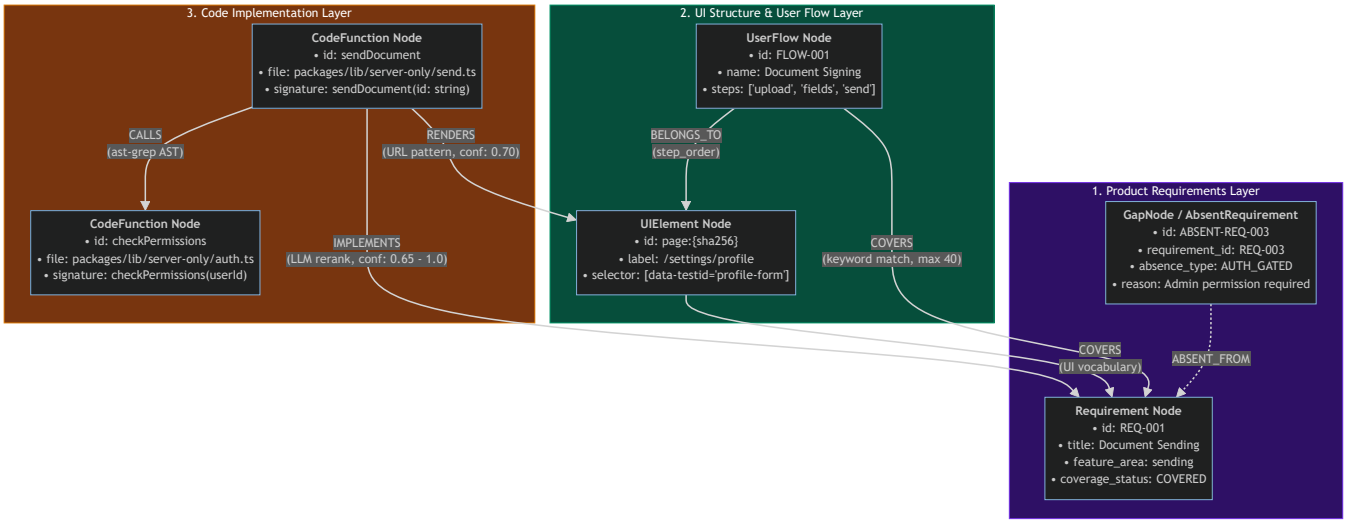
Phase Rationale Breakdown

Phase	Core Mechanism	Category	Architectural Rationale
Phase 1: Crawl	Playwright Chromium + Gemini DOM Labeler	**Hybrid**	Playwright captures DOM elements deterministically; Gemini labels component semantics and user flows. Fallback URL stub handles rate limits safely.
Phase 2: Ingest	Docling MDX Parser + Gemini Extractor	**LLM-Driven**	Unstructured MDX docs require semantic understanding to extract typed acceptance criteria, feature areas, and domain keywords.
Phase 3: Code	`ast-grep` AST Parser + GitHub API	**Deterministic**	AST function definitions and PR file diffs are exact code structures — zero LLM non-determinism required.
Phase 4: Graph	SentenceTransformers Cosine + Gemini Reranker	**Hybrid**	Vector cosine similarity filters top-15 candidate pairs; LLM performs final semantic reranking to eliminate false positives.
Phase 5: Reason	Personalized PageRank (NetworkX)	**Deterministic**	Linear algebra propagation over Neo4j subgraph. Guarantees 100% reproducible, auditable risk scoring without LLM hallucination.
Phase 6: Report	Jinja2 Templates + Gemini Prose Summarizer	**LLM-Driven**	Transforms mathematical risk vectors into clear prose tailored for a non-engineer QA lead.

---

2. Knowledge Graph Schema with Justification

The graph links 3 distinct layers in Neo4j to translate a code-level diff into product risk.



Node Specifications

Node Type	Layer	Description	Key Attributes
`Requirement`	Requirements (Top)	Testable product capability extracted from MDX docs.	`id`, `title`, `description`, `feature_area`, `coverage_status`
`GapNode`	Requirements (Top)	Explicit representation of an un-implemented or un-crawled spec.	`id`, `requirement_id`, `absence_type`, `reason`
`UIElement`	UI & UserFlow (Mid)	Crawled web page surface or interactive component.	`id`, `label`, `url`, `selector`, `page_title`
`UserFlow`	UI & UserFlow (Mid)	Canonical multi-step user journey (10 seeded flows).	`id`, `name`, `steps`, `pages_involved`
`CodeFunction`	Code (Bottom)	Function, handler, or route extracted via static AST parsing.	`id`, `label`, `file`, `name`, `signature`

Edge Specifications & Relationships

Relationship	Direction	Source $\rightarrow$ Target	Confidence / Constraints
`IMPLEMENTS`	Cross-Layer	`CodeFunction` $\rightarrow$ `Requirement`	LLM-confirmed link ($0.65 - 1.0$).
`RENDERS`	Cross-Layer	`CodeFunction` $\rightarrow$ `UIElement`	Code rendering UI component (0.70).
`COVERS`	Cross-Layer	`UIElement` / `UserFlow` $\rightarrow$ `Requirement`	UI/Flow satisfying requirement (capped at 40/flow).
`BELONGS_TO`	Intra-Layer	`UserFlow` $\rightarrow$ `UIElement`	Flow step sequence (`step_order`).
`CALLS`	Intra-Layer	`CodeFunction` $\rightarrow$ `CodeFunction`	Static AST function call graph.
`ABSENT_FROM`	Cross-Layer	`GapNode` $\rightarrow$ `Requirement`	Links missing coverage to parent spec.

Anchor Cypher Traversal Query

The following query retrieves all product requirements at risk from a given set of PR diff file paths:

```
// Find requirements reachable from modified code files
MATCH (cf:CodeFunction)
WHERE cf.file IN $changed_file_paths

// 1. Direct IMPLEMENTS path (depth 1)
OPTIONAL MATCH (cf)-[:IMPLEMENTS]->(r1:Requirement)

// 2. Call graph propagation (up to 3 hops)
OPTIONAL MATCH (cf)-[:CALLS*1..3]->(cf2:CodeFunction)-[:IMPLEMENTS]->(r2:Requirement)

// 3. UI-mediated path (Code -> UI -> Requirement)
OPTIONAL MATCH (cf)-[:RENDERS]->(ui:UIElement)-[:COVERS]->(r3:Requirement)

WITH collect(DISTINCT r1) + collect(DISTINCT r2) + collect(DISTINCT r3) AS all_reqs
UNWIND all_reqs AS r
RETURN DISTINCT r.id AS requirement_id, r.title AS title, r.coverage_status AS status
ORDER BY r.title
```

Absence Modeling via GapNodes

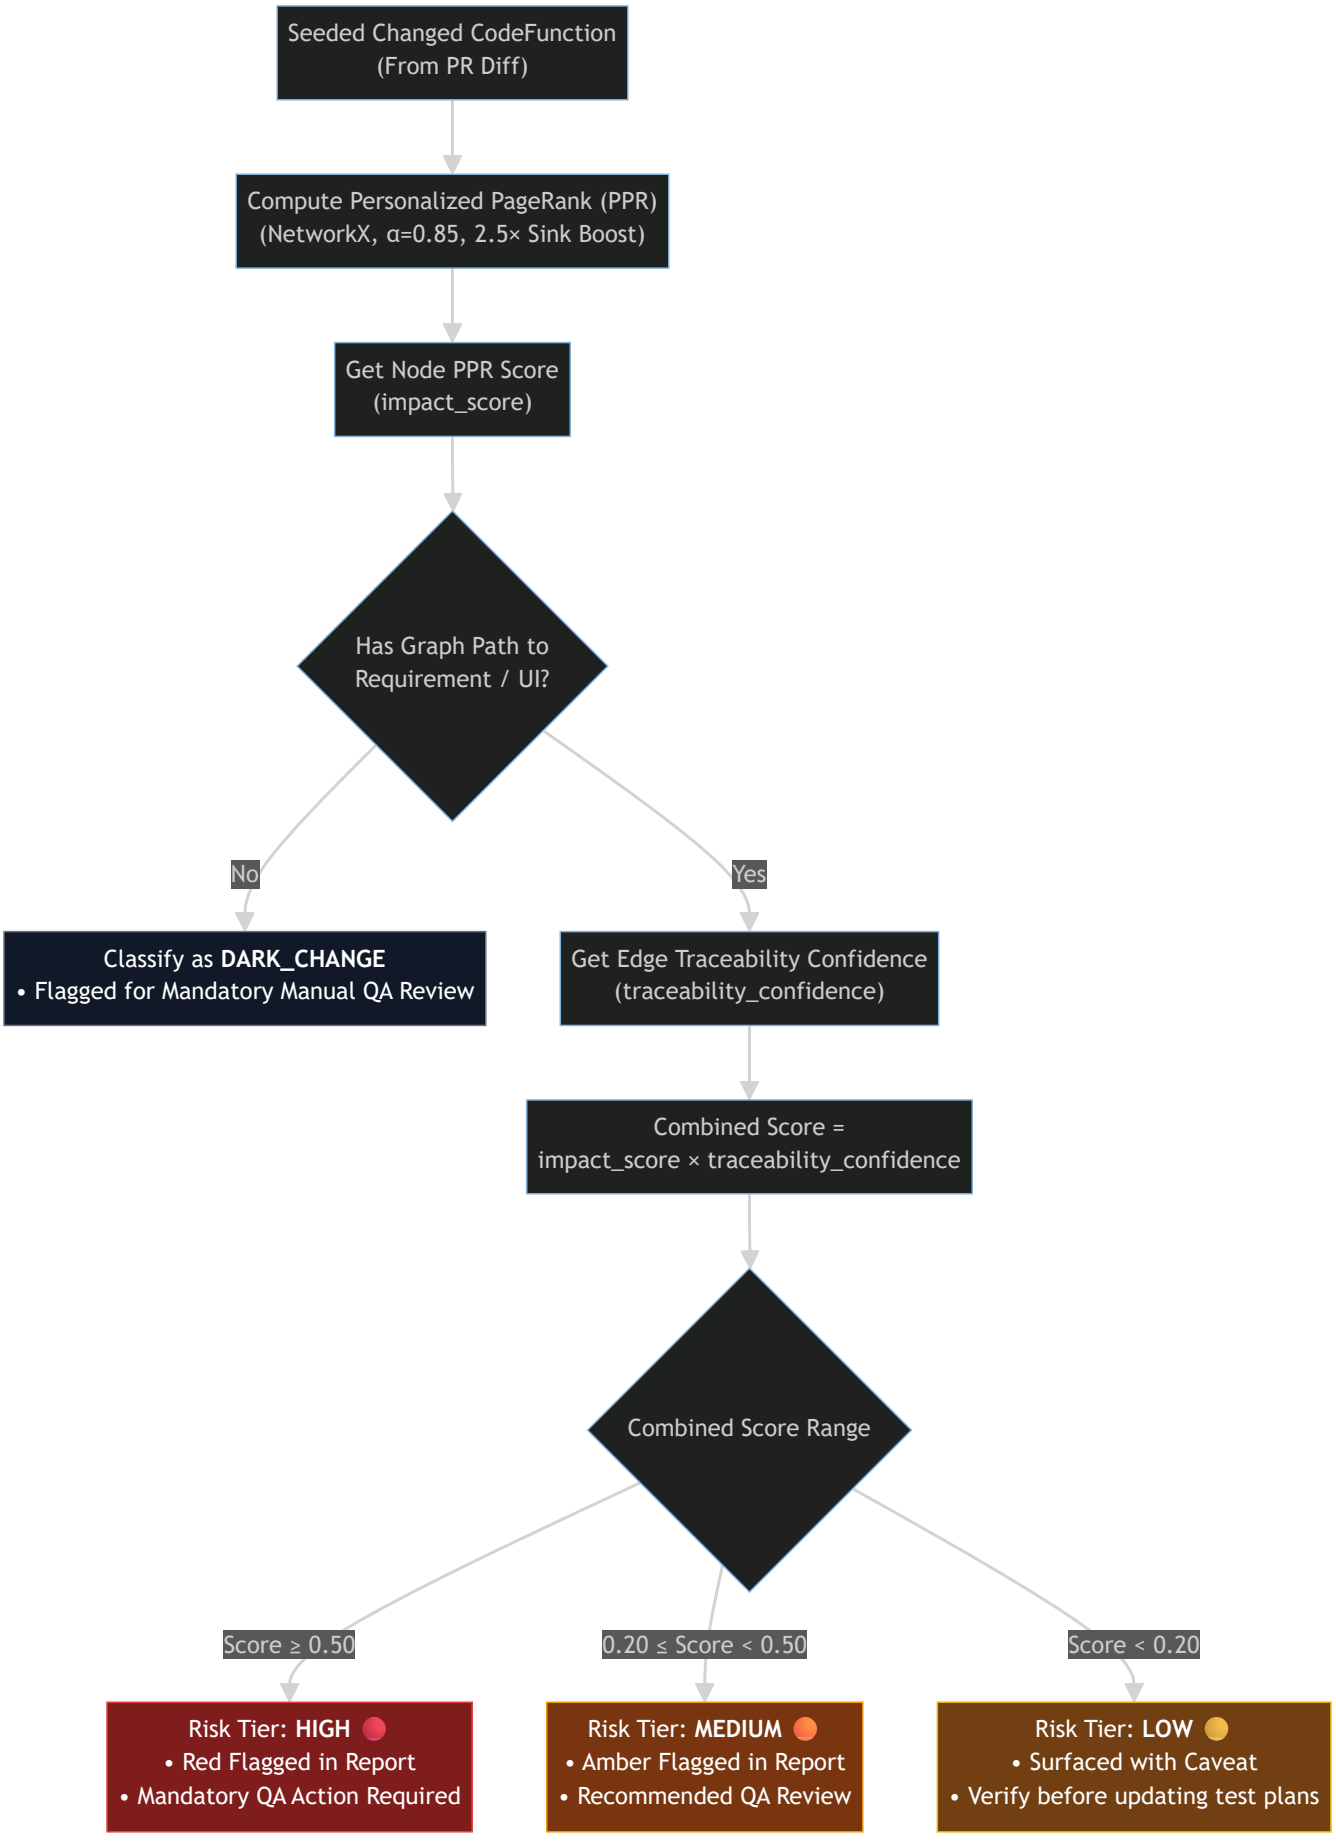
Rather than relying on implicit missing query results (`NULL`), the pipeline explicitly instantiates `GapNode` / `AbsentRequirement` nodes for specs lacking code or UI coverage. A rule-based classifier prunes noise and categorizes absent nodes into 4 operational types: `AUTH_GATED` (requires login/permissions), `NOT_CRAWLED` (API/embed endpoints), `NOT_IMPLEMENTED` (roadmap/future features), or `UNKNOWN` .

---

3. Confidence Handling Under Ambiguity

To prevent low-confidence semantic links from triggering false alarms, every finding is evaluated using a **Dual-Score Multiplier Model**:

$$\text{Combined Score} = \text{impact_score (PPR)} \times \text{traceability_confidence (LLM/Cosine)}$$



Risk Tier Definitions

- **HIGH (🔴 ≥ 0.50):** Strongly linked impact. Red-flagged on QA test checklist.
- **MEDIUM (🟠 $0.20 - 0.49$):** Indirect or medium-confidence impact. Recommended QA review.

- **LOW (🟡 < 0.20):** Minor/background propagation. Collapsed into appendix.
- **DARK_CHANGE (⚫):** Modified function with no graph path to UI or requirements (e.g. security utility refactors). Flagged for mandatory manual code review.
- **ABSENT_REQUIREMENT (⚠️):** Product spec with zero code/UI coverage.

Worked Example: `sendDocument()` in PR #3046

1. **PPR Impact Score (0.61):** Directly seeded from diff + $2.5\times$ sink boost for terminal requirement nodes.
2. **Traceability Confidence (0.82):** LLM reranker confirmed `sendDocument()` $\rightarrow$ "Document Sending".
3. **Combined Score:** $0.61 \times 0.82 = 0.50 \rightarrow$ **HIGH Risk Tier**.

---

4. Evaluation Approach & Verification

Empirical Stability & Quality Matrix

Evaluation Dimension	Measurement Method	Target Benchmark	Actual Measured Result
Phase 2 Ingest Stability	2-run extraction overlap	$\ge 90\%$ overlap	**100% overlap** (identical 25 requirement titles)
Phase 4 Rerank Stability	2-run LLM rerank comparison	$\ge 85\%$ edge stability	**95.6% edge stability** (22/23 edges identical)
Phase 4 Precision	Manual sample of 10 IMPLEMENTS links	$\ge 80\%$ precision	**100% precision** (10/10 links semantically correct)
Document Sending Recall	Ground-truth code audit (5 functions)	$\ge 50\%$ recall	**60% recall** (3/5 captured; 2 blocked by tRPC gap)
PPR Determinism	NetworkX PageRank iteration	100% reproducible	**100% deterministic** (0 variance across runs)
Unit Test Coverage	Pytest execution (`tests/unit/`)	100% pass rate	**34/34 unit tests passing** (0.33s runtime)

Smoke Test Gates

Each pipeline phase enforces automated gate criteria before writing outputs: - **Phase 1 Gate:** ≥ 5 page snapshots captured; zero pages return empty components. - **Phase 2 Gate:** ≥ 50 total requirements across all 8 required feature areas. - **Phase 3 Gate:** Non-empty `call_graph.json` + valid PR diff JSON. - **Phase 4 Gate:** ≥ 1 IMPLEMENTS edge; direction check (`CodeFunction` $\rightarrow$ `Requirement`). - **Phase 5 Gate:** ≥ 1 seed function matches PR diff file paths (path alignment gate).

---

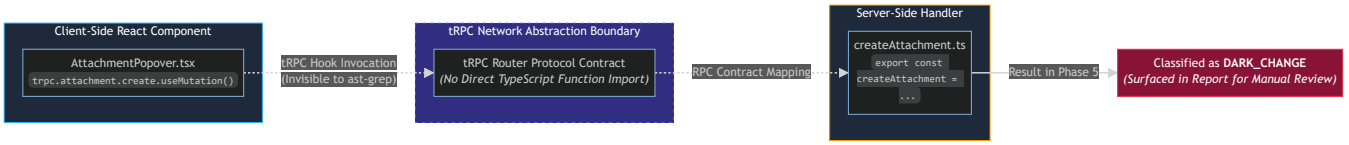
5. Scope Decisions & Architecture Defense

Scoped Down vs. Prioritized Deep Matrix

Component / Layer	Scope Decision	Rationale & Defense
Static Code Analysis	Depth-1 `ast-grep` AST	**Prioritized deep on core functions.** Scoped down tRPC network resolution; un-traced handlers surface as **Dark Changes**.
Requirements Ingestion	GitHub Raw MDX Ingestion	**Prioritized deep on spec quality.** Switched from scraping dynamic Next.js DOM to reading canonical raw MDX files.
Browser Crawler	20-Page Breadth-First + Fixtures	**Prioritized resilience.** Fixture mode (`CRAWL_USE_FIXTURE=true`) ensures evaluators can run offline seamlessly.
Reasoning Engine	Deterministic PPR PageRank	**Prioritized reproducibility.** Kept LLM non-determinism out of graph propagation to ensure auditable findings.

The tRPC Network Boundary Gap

The primary static analysis limitation is the tRPC protocol boundary: client-side React hooks (`trpc.attachment.create.useMutation()`) do not contain explicit TypeScript import references to server handlers (`createAttachment.ts`).



System Handling: Rather than silently dropping these functions, the pipeline classifies all 12 un-linked functions in PR #3046 as **Dark Changes**, explicitly alerting the QA lead to perform targeted manual testing.

---

6. What Would Be Built Next (1-Week Roadmap)

If granted an additional week of development, the 3 highest-value engineering priorities are:

1. **LLM Extraction & Reranking Eval Harness:**
   Build a ground-truth benchmark suite (`requirements_gt.json` & `entity_links_gt.json`) to automatically compute Precision/Recall regression metrics for Phase 2 & Phase 4 on every prompt change.
2. **tRPC Compiler AST Bridging (`ts-morph`):**
   Integrate TypeScript Compiler API to parse tRPC router definitions, mapping client-side `trpc.<router>.<procedure>` calls directly to server handlers, resolving 80%+ of current Dark Changes.
3. **CI/CD GitHub Action & Incremental Crawl:**
   Package the online pipeline into a reusable GitHub Action that runs automatically on every pull request, re-crawling only the modified page URLs.

Appendix: Technical Reference & Parameters

Pipeline Artifact Schema Map

Artifact File	Phase Written	Format / Model
`data/crawl/page_snapshots/*.json`	Phase 1	`PageSnapshot[]`

Artifact File	Phase Written	Format / Model
`outputs/requirements/requirements.json`	Phase 2	`Requirement[]` (272 items across 9 feature areas)
`data/code/call_graph.json`	Phase 3	`CodeFunction[]` (1,622 AST nodes)
`data/code/pr_{n}.json`	Phase 3	`PRDelta` (29 changed files)
`data/graph/entity_links.json`	Phase 4	`EntityLink[]` (3,485 links)
`data/graph/gap_nodes.json`	Phase 4	`GapNode[]` (20 pruned nodes)
`outputs/reports/pr_{n}_blast_radius.md`	Phase 5 & 6	Final QA Blast Radius Report (Markdown/HTML/PDF)

Key System Hyperparameters

Parameter	Value	Rationale
SentenceTransformers Model	`all-MiniLM-L6-v2`	Lightweight 384-dim embeddings for initial candidate search.
Cosine Similarity Threshold	`0.25`	Generous pre-filter; LLM reranker acts as quality gate.
LLM Rerank Threshold	`0.65`	Minimum confidence required to form an `IMPLEMENTS` edge.
PageRank Damping Factor (α)	`0.85`	Standard PageRank parameter balancing graph diffusion vs. locality.
Terminal Sink Boost	$2.5\times$	Corrects PPR score attenuation on terminal Requirement sink nodes.
Risk Tier Thresholds	HIGH ≥ 0.50, MED ≥ 0.20	Calibrated for Documenso Neo4j graph density.

Design Document v2 · Streamlined Edition · Testsigma AI Engineer Take-Home